




Google Developer Group
Chandigarh University

Gemini Study Jams

Session 2 - Supervised PEFT using LoRA

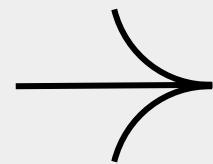
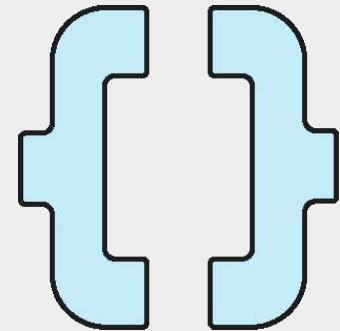


What will we do today?

1. The main loop of the session will be:


Code -> Theory -> Code

2. We will work through the given code and stop and understand the theory behind it at **Exhibitions**.
3. It is **recommended** that you have the file “Session 2 - PEFT with LoRA.ipynb” open.



E1: LoRA decreases complexity!

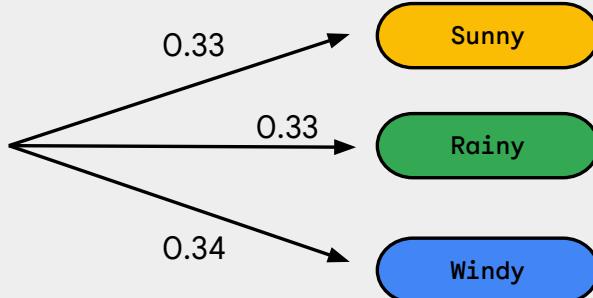
1. During full fine-tuning, **the model is initialized to pre-trained weights Φ and updated to $\Phi + \Delta\Phi$** by repeatedly following the gradient to maximize the conditional language modeling objective.
2. One of the main drawbacks for full fine-tuning is that **for each downstream task, we learn a different set of parameters $\Delta\Phi$** whose dimension $|\Delta\Phi|$ equals $|\Phi|$.
3. This means that for each downstream task, for a model like GPT3, **$|\Phi| \approx 175 \text{ Billion}$** , storing and deploying many independent instances of fine-tuned models can be challenging, if at all feasible.



Overly Simplified Visualization : A Sunny Model

Pre-Trained Model


Today is a _____ day.



$$\Phi = [0.33, 0.33, 0.34]$$

Fine Tuned Model

Today is a _____ day.




$$\Phi = [0.6, 0.2, 0.2]$$

$$\Delta\Phi = [0.27, -0.13, -0.14]$$

Now imagine 175 billion parameter updates for GPT3 - unique to each task!

E1: LoRA decreases complexity!

1. In LoRA, the task-specific parameter increment $\Delta\Phi = \Delta\Phi(\Theta)$ is **further encoded by a much smaller-sized set of parameters** Θ with $|\Theta| \ll |\Phi|$.
2. When the pre-trained model is GPT-3 175B, **the number of trainable parameters $|\Theta|$ can be as small as 0.01% of $|\Phi|$** .




Now back to Code!

E2: How does LoRA work?

1. Consider a single Transformer layer, single attention head for clarity.
2. You start with an input sequence:

$X \in \mathbb{R}$ where X has dimensions $T * d$ where

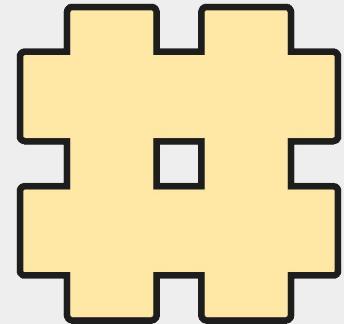
- T : sequence length
 - d : embedding dimension
3. Each token is a row vector in X .



E2: How does LoRA work?

1. The attention mechanism does **not** operate directly on X .
2. Instead, it projects X into three different representation spaces:
 - **Query (Q)** — “What am I looking for?”
 - **Key (K)** — “What do I contain?”
 - **Value (V)** — “What information do I give if selected?”
3. And then attention is computed using these three.
4. Mathematically:

$$Q = XW(Q) \mid V = XW(V) \mid K = XW(K)$$



E2: How does LoRA work?

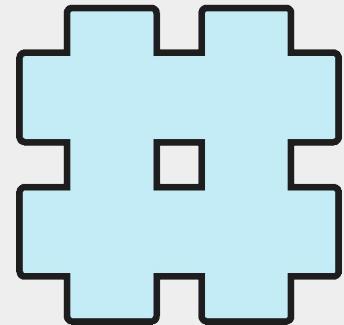
1. During Fine Tuning, the projection matrices are learned as follows:

$$Q = X(W(Q) + \Delta W(Q))$$

2. LoRA **does not change** the attention mechanism.
3. It changes how the projection matrices are *learned*. LoRA enforces:

$$\Delta W(Q) = BA$$

Where A has dims $r * d$ and B has dimension $d * r$ where $r \ll d$. The original weight matrix $W(Q)$ is frozen, and only A and B are trained.



E3: LoRA Evaluation

(a) Language modeling


- Perplexity (PPL)
- Negative log-likelihood (NLL)

(b) Classification / structured prediction

- Accuracy
- Precision / Recall / F1
- ROC-AUC

(c) Generation tasks

- BLEU, ROUGE, METEOR (summarization, translation)
- Exact Match (EM), F1 (QA)
- Human evaluation (helpfulness, factuality, style adherence)



E4: Considerations

Choose target layers carefully (Q, V usually most effective)

Select rank r wisely: low r captures most gains; high r has diminishing returns

Tune scaling α to control update magnitude and stability

Initialize properly (base weights frozen; LoRA starts as no-op)


Use higher learning rates; avoid heavy regularization

Best for domain/style adaptation; limited for learning new capabilities

Check in-domain vs out-of-domain generalization

Decide whether to merge adapters at inference

Monitor memory, speed, and perplexity together



Quiz Time!

Doubts?

Ask away!



Google Developer Group
Chandigarh University

Thank You!

Fine Tuning Basics & Supervised Fine Tuning - End

